

Car_Price_Prediction

August 29, 2024

```
[177]: import pandas as pd
import numpy as np
from scipy import stats
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
pd.options.display.float_format='{:, .2f}'.format # to round all number to 0.2
```

0.0.1 About Dataset

The price of a car depends on a lot of factors like the goodwill of the brand of the car, features of the car, horsepower and the mileage it gives, and many more. Car price prediction is one of the major research areas in machine learning. This dataset can be used to train a car price prediction model.

0.0.2 Dataset Columns

Column Name	Description
Year	The year the car was manufactured.
Selling_Price	The price at which the car is being sold.
Present_Price	The current ex-showroom price of the car.
Driven_kms	The total kilometers the car has been driven.
Fuel_Type	The type of fuel the car uses (e.g., Petrol, Diesel).
Selling_type	The type of seller (e.g., Dealer, Individual).
Transmission	The type of transmission (e.g., Manual, Automatic).
Owner	The number of previous owners the car has had.
No_of_Years	The number of years since the car was manufactured.

```
[178]: data=pd.read_csv('Car_Price.csv',)
data
```

```
[178]:   Car_Name  Year  Selling_Price  Present_Price  Driven_kms  Fuel_Type  \
0      ritz  2014           3.35           5.59       27000     Petrol
1       sx4  2013           4.75           9.54       43000     Diesel
2      ciaz  2017           7.25           9.85        6900     Petrol
3  wagon r  2011           2.85           4.15        5200     Petrol
4     swift  2014           4.60           6.87       42450     Diesel
```

```

..      ...
296    city  2016          9.50          11.60          33988    Diesel
297    brio  2015          4.00           5.90          60000    Petrol
298    city  2009          3.35          11.00          87934    Petrol
299    city  2017         11.50          12.50           9000    Diesel
300    brio  2016          5.30           5.90          5464    Petrol

```

```

      Selling_type Transmission  Owner
0          Dealer      Manual      0
1          Dealer      Manual      0
2          Dealer      Manual      0
3          Dealer      Manual      0
4          Dealer      Manual      0
..      ...
296    Dealer      Manual      0
297    Dealer      Manual      0
298    Dealer      Manual      0
299    Dealer      Manual      0
300    Dealer      Manual      0

```

[301 rows x 9 columns]

1 Data_Size

```
[179]: print('The Size Of Data Frame is :',data.shape)
```

The Size Of Data Frame is : (301, 9)

```
[180]: data.columns
```

```
[180]: Index(['Car_Name', 'Year', 'Selling_Price', 'Present_Price', 'Driven_kms',
          'Fuel_Type', 'Selling_type', 'Transmission', 'Owner'],
          dtype='object')
```

2 Data_Types

```
[181]: data.dtypes
```

```
[181]: Car_Name      object
      Year         int64
      Selling_Price  float64
      Present_Price  float64
      Driven_kms     int64
      Fuel_Type      object
      Selling_type   object
      Transmission   object
```

```
Owner          int64
dtype: object
```

```
[182]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Car_Name        301 non-null   object
1   Year            301 non-null   int64
2   Selling_Price   301 non-null   float64
3   Present_Price   301 non-null   float64
4   Driven_kms      301 non-null   int64
5   Fuel_Type       301 non-null   object
6   Selling_type    301 non-null   object
7   Transmission    301 non-null   object
8   Owner           301 non-null   int64
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB
```

3 Missing__Values

```
[183]: data.isna().sum()
```

```
[183]: Car_Name        0
      Year           0
      Selling_Price   0
      Present_Price   0
      Driven_kms      0
      Fuel_Type       0
      Selling_type    0
      Transmission    0
      Owner           0
      dtype: int64
```

```
[184]: data.isnull().mean()*100
```

```
[184]: Car_Name        0.00
      Year           0.00
      Selling_Price   0.00
      Present_Price   0.00
      Driven_kms      0.00
      Fuel_Type       0.00
      Selling_type    0.00
      Transmission    0.00
```

```
Owner          0.00
dtype: float64
```

4 Duplicated Values

```
[185]: data.duplicated().sum()
```

```
[185]: 2
```

```
[186]: data.drop_duplicates(inplace=True)
data.duplicated().sum()
```

```
[186]: 0
```

```
[187]: print('The Size Of Data Frame after clean : ',data.shape)
```

```
The Size Of Data Frame after clean : (299, 9)
```

5 Unique Values

```
[188]: for col in data.columns:
        print(f'{col} number of unique values is : {data[col].nunique()} \n↳ {data[col].unique()}')
        print('*****')
```

```
Car_Name number of unique values is : 98
['ritz' 'sx4' 'ciaz' 'wagon r' 'swift' 'vitara brezza' 's cross'
'alto 800' 'ertiga' 'dzire' 'alto k10' 'ignis' '800' 'baleno' 'omni'
'fortuner' 'innova' 'corolla altis' 'etios cross' 'etios g' 'etios liva'
'corolla' 'etios gd' 'camry' 'land cruiser' 'Royal Enfield Thunder 500'
'UM Renegade Mojave' 'KTM RC200' 'Bajaj Dominar 400'
'Royal Enfield Classic 350' 'KTM RC390' 'Hyosung GT250R'
'Royal Enfield Thunder 350' 'KTM 390 Duke ' 'Mahindra Mojo XT300'
'Bajaj Pulsar RS200' 'Royal Enfield Bullet 350'
'Royal Enfield Classic 500' 'Bajaj Avenger 220' 'Bajaj Avenger 150'
'Honda CB Hornet 160R' 'Yamaha FZ S V 2.0' 'Yamaha FZ 16'
'TVS Apache RTR 160' 'Bajaj Pulsar 150' 'Honda CBR 150' 'Hero Extreme'
'Bajaj Avenger 220 dtsi' 'Bajaj Avenger 150 street' 'Yamaha FZ v 2.0'
'Bajaj Pulsar NS 200' 'Bajaj Pulsar 220 F' 'TVS Apache RTR 180'
'Hero Passion X pro' 'Bajaj Pulsar NS 200' 'Yamaha Fazer '
'Honda Activa 4G' 'TVS Sport ' 'Honda Dream Yuga '
'Bajaj Avenger Street 220' 'Hero Splender iSmart' 'Activa 3g'
'Hero Passion Pro' 'Honda CB Trigger' 'Yamaha FZ S '
'Bajaj Pulsar 135 LS' 'Activa 4g' 'Honda CB Unicorn'
'Hero Honda CBZ extreme' 'Honda Karizma' 'Honda Activa 125' 'TVS Jupyter'
'Hero Honda Passion Pro' 'Hero Splender Plus' 'Honda CB Shine'
'Bajaj Discover 100' 'Suzuki Access 125' 'TVS Wego' 'Honda CB twister'
```

```
'Hero Glamour' 'Hero Super Splendor' 'Bajaj Discover 125' 'Hero Hunk'
'Hero Ignitor Disc' 'Hero CBZ Xtreme' 'Bajaj ct 100' 'i20' 'grand i10'
'i10' 'eon' 'xcent' 'elantra' 'creta' 'verna' 'city' 'brio' 'amaze'
'jazz']
```

Year number of unique values is : 16

```
[2014 2013 2017 2011 2018 2015 2016 2009 2010 2012 2003 2008 2006 2005
2004 2007]
```

Selling_Price number of unique values is : 156

```
[ 3.35  4.75  7.25  2.85  4.6   9.25  6.75  6.5   8.75  7.45  6.85  7.5
 6.1   2.25  7.75  3.25  2.65  4.9   4.4   2.5   2.9   3.    4.15  6.
 1.95  3.1   2.35  4.95  5.5   2.95  4.65  0.35  5.85  2.55  1.25  1.05
 5.8   14.9  23.   18.   16.   2.75  3.6   4.5   4.1  19.99  6.95 18.75
23.5  33.   19.75  4.35 14.25  3.95  1.5   5.25 14.5  14.73 12.5  3.49
35.   5.9   3.45  3.8  11.25  3.51  4.   20.75 17.   7.05  9.65  1.75
 1.7   1.65  1.45  1.35  1.2   1.15  1.11  1.1   1.    0.95  0.9   0.75
 0.8   0.78  0.72  0.65  0.6   0.55  0.52  0.51  0.5   0.48  0.45  0.42
 0.4   0.38  0.31  0.3   0.27  0.25  0.2   0.18  0.17  0.16  0.15  0.12
 0.1   5.75  5.15  7.9   4.85 11.75  3.15  6.45  3.5   8.25  5.11  2.7
 6.15 11.45  3.9   9.1   4.8   2.    5.35  6.25  5.95  5.2   3.75 12.9
 5.    5.4   7.2  10.25  8.5   8.4   9.15  6.6   3.65  8.35  6.7   5.3
10.9  8.65  9.7   2.1   8.99  7.4   5.65 10.11  6.4   8.55  9.5  11.5 ]
```

Present_Price number of unique values is : 148

```
[ 5.59  9.54  9.85  4.15  6.87  9.83  8.12  8.61  8.89  8.92
 3.6   10.38  9.94  7.71  7.21 10.79  5.09  7.98  3.95  5.71
 8.01  3.46  4.41  4.99  5.87  6.49  5.98  4.89  7.49  9.95
 8.06  7.74  7.2   2.28  3.76  7.87  3.98  7.15  2.69 12.04
 9.29 30.61 19.77 10.21 15.04  7.27 18.54  6.8  35.96 18.61
 7.7  36.23  6.95 23.15 20.45 13.74 20.91  6.76 12.48  8.93
14.68 12.35 22.83 14.89  7.85 25.39 13.46 23.73 92.6  6.05
16.09 13.7  22.78 18.64  1.9   1.82  1.78  1.6   1.47  2.37
 3.45  1.5   2.4   1.4   1.26  1.17  1.75  0.95  0.8   0.87
 0.84  0.82  0.81  0.74  1.2   0.787  0.99  0.94  0.826  0.55
 0.88  0.51  0.52  0.54  0.73  0.83  0.64  0.72  1.05  0.57
 0.48  0.58  0.47  0.75  0.65  0.32  6.79  5.7   4.6   4.43
 7.13  8.1  14.79 13.6   9.4   8.4   5.43  7.6   9.9   6.82
 5.35  7.    5.97  5.8   8.7  10.   7.5   5.9  14.   11.8
 8.5   7.9   6.4   6.1  13.09 11.6  11.   12.5 ]
```

Driven_kms number of unique values is : 206

```
[ 27000 43000  6900  5200 42450  2071 18796 33429 20273 42367
 2135 51000 15000 26000 77427 41678 35500 41442 25000 2400
50000 45280 56879 20000 55138 16200 44542 45000 51439 54200
39000 49998 48767 127000 10079 62000 24524 46706 58000 45780
64532 65000 25870 37000 104707 40000 135000 90000 70000 40534
39485 41000 40001 40588 78000 47000  6000 11000 59000 88000
```

```

12000 71000 56001 83000 36000 72000 135154 80000 89000 23000
38000 197176 142000 56000 58242 75000 29000 8700 50024 3000
1400 4000 1200 4100 21700 16500 18000 7000 35000 17000
17500 33000 14000 5400 5700 46500 11500 1300 5000 3500
500 11800 23500 16000 16600 32000 19000 24000 31000 13000
8000 4300 8600 14500 1000 42000 5500 6700 13700 38600
30000 213000 60000 21000 1900 22000 55000 49000 500000 53000
92233 28200 53460 28282 3493 12479 34797 3435 21125 35775
43535 22671 31604 20114 36100 12500 45078 38488 77632 61381
36198 22517 24678 57000 52132 15001 12900 4492 15141 11849
68000 60241 23709 32322 35866 34000 35934 56701 31427 48000
54242 53675 49562 40324 36054 29223 5600 40023 16002 40026
21200 19434 18828 69341 69562 27600 61203 30753 24800 21780
40126 14465 50456 63000 9010 9800 15059 28569 44000 10980
33019 60076 33988 87934 9000 5464]

```

```
*****
```

```
Fuel_Type number of unique values is : 3
```

```
['Petrol' 'Diesel' 'CNG']
```

```
*****
```

```
Selling_type number of unique values is : 2
```

```
['Dealer' 'Individual']
```

```
*****
```

```
Transmission number of unique values is : 2
```

```
['Manual' 'Automatic']
```

```
*****
```

```
Owner number of unique values is : 3
```

```
[0 1 3]
```

```
*****
```

```
[ ]: data.columns.str.strip()
```

```
[189]: data['Owner'] = data['Owner'].replace(to_replace=3, value=2)
print("'Owner' variable has {} unique categories: {}".format(data['Owner'].
↪nunique(), data['Owner'].unique()))
```

```
'Owner' variable has 3 unique categories: [0 1 2]
```

```
[190]: data['Current_Year'] = 2024
data['No_of_Years'] = data['Current_Year'] - data['Year']
data.head()
```

```
[190]: Car_Name  Year  Selling_Price  Present_Price  Driven_kms  Fuel_Type  \
0      ritz    2014           3.35           5.59         27000     Petrol
1       sx4    2013           4.75           9.54         43000     Diesel
2      ciaz    2017           7.25           9.85          6900     Petrol
3  wagon r    2011           2.85           4.15          5200     Petrol
4    swift    2014           4.60           6.87         42450     Diesel
```

	Selling_type	Transmission	Owner	Current_Year	No_of_Years
0	Dealer	Manual	0	2024	10
1	Dealer	Manual	0	2024	11
2	Dealer	Manual	0	2024	7
3	Dealer	Manual	0	2024	13
4	Dealer	Manual	0	2024	10

```
[191]: data.drop(['Current_Year', 'Car_Name'], inplace=True, axis=1)
```

6 Data_Preview

```
[192]: data.sample(2)
```

```
[192]:
```

	Year	Selling_Price	Present_Price	Driven_kms	Fuel_Type	Selling_type	\
125	2009	0.90	1.75	40000	Petrol	Individual	
80	2016	14.73	14.89	23000	Diesel	Dealer	

	Transmission	Owner	No_of_Years
125	Manual	0	15
80	Manual	0	8

```
[193]: data.head()
```

```
[193]:
```

	Year	Selling_Price	Present_Price	Driven_kms	Fuel_Type	Selling_type	\
0	2014	3.35	5.59	27000	Petrol	Dealer	
1	2013	4.75	9.54	43000	Diesel	Dealer	
2	2017	7.25	9.85	6900	Petrol	Dealer	
3	2011	2.85	4.15	5200	Petrol	Dealer	
4	2014	4.60	6.87	42450	Diesel	Dealer	

	Transmission	Owner	No_of_Years
0	Manual	0	10
1	Manual	0	11
2	Manual	0	7
3	Manual	0	13
4	Manual	0	10

```
[194]: data.tail()
```

```
[194]:
```

	Year	Selling_Price	Present_Price	Driven_kms	Fuel_Type	Selling_type	\
296	2016	9.50	11.60	33988	Diesel	Dealer	
297	2015	4.00	5.90	60000	Petrol	Dealer	
298	2009	3.35	11.00	87934	Petrol	Dealer	
299	2017	11.50	12.50	9000	Diesel	Dealer	
300	2016	5.30	5.90	5464	Petrol	Dealer	

	Transmission	Owner	No_of_Years
--	--------------	-------	-------------

296	Manual	0	8
297	Manual	0	9
298	Manual	0	15
299	Manual	0	7
300	Manual	0	8

7 Statistical_OverView

```
[195]: data.describe().T
```

```
[195]:
```

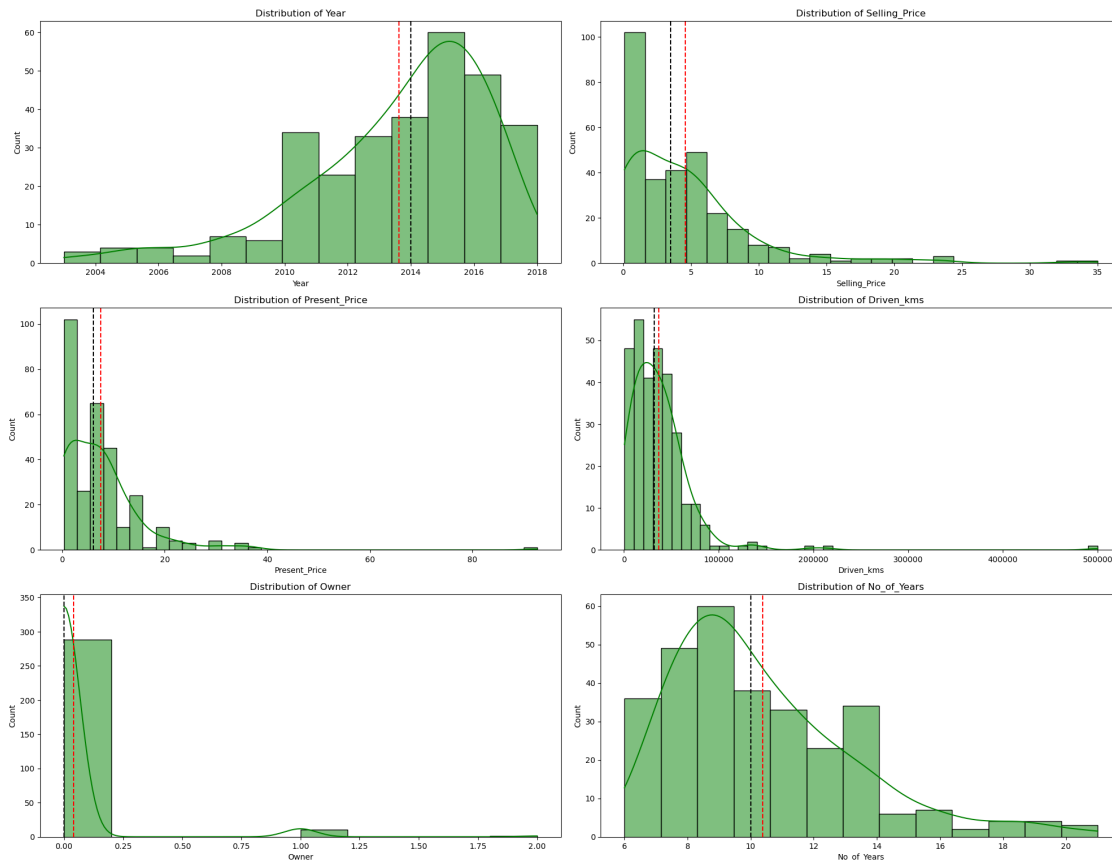
	count	mean	std	min	25%	50% \
Year	299.00	2,013.62	2.90	2,003.00	2,012.00	2,014.00
Selling_Price	299.00	4.59	4.98	0.10	0.85	3.51
Present_Price	299.00	7.54	8.57	0.32	1.20	6.10
Driven_kms	299.00	36,916.75	39,015.17	500.00	15,000.00	32,000.00
Owner	299.00	0.04	0.21	0.00	0.00	0.00
No_of_Years	299.00	10.38	2.90	6.00	8.00	10.00

	75%	max
Year	2,016.00	2,018.00
Selling_Price	6.00	35.00
Present_Price	9.84	92.60
Driven_kms	48,883.50	500,000.00
Owner	0.00	2.00
No_of_Years	12.00	21.00

8 Univariate analysis

1-Numerical Data (Histogram/Distplot - Box Plot - Summary Statistics)

```
[196]: numeric_columns = data.select_dtypes('number')
plt.figure(figsize=(20,20))
for i ,e in enumerate(numeric_columns):
    plt.subplot(4,2,i+1)
    sns.histplot(data[e],kde=True,color='g')
    plt.axvline(data[e].mean(), color='r', linestyle='--')
    plt.axvline(data[e].median(), color='black', linestyle='--')
    plt.title('Distribution of '+e)
plt.tight_layout()
```

mean > median (right-skewed)

```
[197]: data.select_dtypes('number').describe()
```

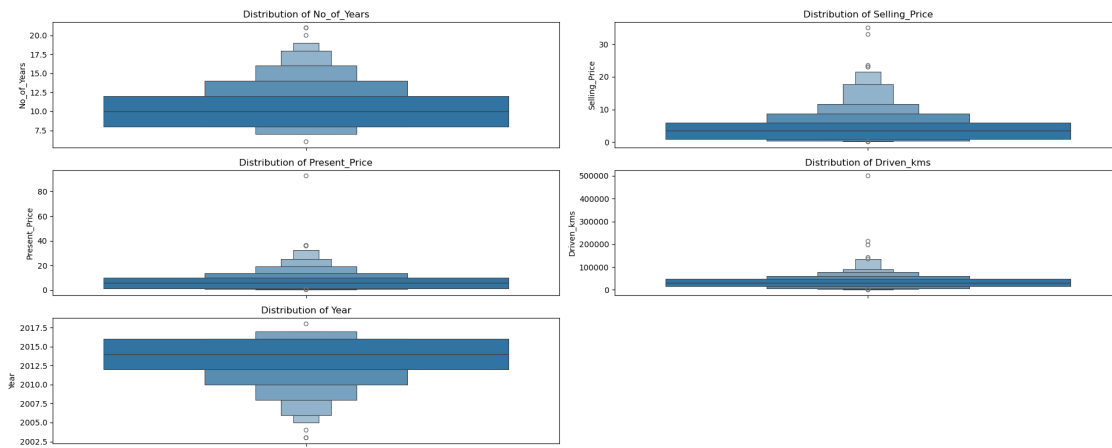
```
[197]:
```

	Year	Selling_Price	Present_Price	Driven_kms	Owner	No_of_Years
count	299.00	299.00	299.00	299.00	299.00	299.00
mean	2,013.62	4.59	7.54	36,916.75	0.04	10.38
std	2.90	4.98	8.57	39,015.17	0.21	2.90
min	2,003.00	0.10	0.32	500.00	0.00	6.00
25%	2,012.00	0.85	1.20	15,000.00	0.00	8.00
50%	2,014.00	3.51	6.10	32,000.00	0.00	10.00
75%	2,016.00	6.00	9.84	48,883.50	0.00	12.00
max	2,018.00	35.00	92.60	500,000.00	2.00	21.00

note that mean > median in table and graph

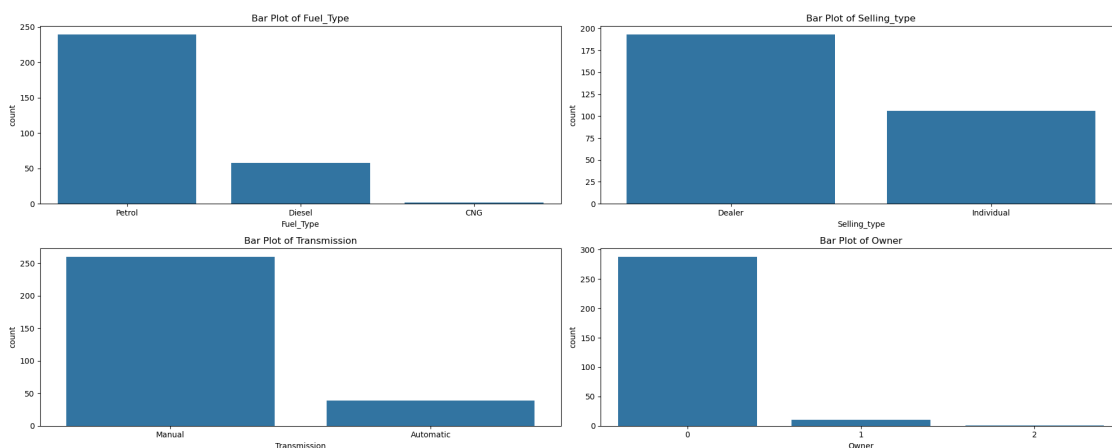
```
[198]: columns_to_plot = ['No_of_Years', 'Selling_Price', 'Present_Price', 'Driven_kms', 'Year']
plt.figure(figsize=(20,8))
for i,col in enumerate(columns_to_plot):
    plt.subplot(3,2,i+1)
```

```
sns.boxenplot(data[col])
plt.title('Distribution of '+col)
plt.tight_layout()
```



2- Categorical Data (Bar Plot - Pie Chart)

```
[201]: categorical_columns = ['Fuel_Type', 'Selling_type', 'Transmission', 'Owner']
plt.figure(figsize=(20,8))
for i,col in enumerate(categorical_columns):
    plt.subplot(2,2,i+1)
    sns.countplot(x=data[col])
    plt.title(f'Bar Plot of {col}')
    plt.tight_layout()
```



```
[ ]: categorical_columns = ['Fuel_Type', 'Selling_type', 'Transmission', 'Owner']
plt.figure(figsize=(20,8))
for i ,col in enumerate(categorical_columns):
    plt.subplot(2,2,i+1)
    v_c=data[col].value_counts()
    plt.pie(v_c,autopct='%1.1f%%', startangle=90)
    plt.title(f'Pie Chart of {col}')
    plt.tight_layout()
```

9 Bivariate Analysis

1- Numerical vs. Numerical (Correlation Matrix - Scatter Plot - Scatter Plot With a Trend Line)

```
[202]: correlation_matrix=numeric_columns.corr()
correlation_matrix
```

```
[202]:
```

	Year	Selling_Price	Present_Price	Driven_kms	Owner	\
Year	1.00	0.23	-0.05	-0.53	-0.17	
Selling_Price	0.23	1.00	0.88	0.03	-0.10	
Present_Price	-0.05	0.88	1.00	0.21	-0.02	
Driven_kms	-0.53	0.03	0.21	1.00	0.06	
Owner	-0.17	-0.10	-0.02	0.06	1.00	
No_of_Years	-1.00	-0.23	0.05	0.53	0.17	

	No_of_Years
Year	-1.00
Selling_Price	-0.23
Present_Price	0.05
Driven_kms	0.53
Owner	0.17
No_of_Years	1.00

```
[203]: #regplot as scatter but with line
sns.regplot(x='Present_Price', y='Selling_Price', data=data,scatter_kws={'s':
↵50}, line_kws={'color':'red'})
plt.title('Scatter Plot between Present_Price and Selling_Price')
plt.xlabel('Present_Price')
plt.ylabel('Selling_Price')
plt.show()
```




Strong Correlation

```
[204]: # grouped_data = data.groupby('No_of_Years', as_index=False)['Selling_Price'].
        #         ↳mean().round(2)
        # fig = px.scatter(grouped_data, x='No_of_Years', y='Selling_Price',
        #                 title='Scatter plot of No of Years vs Selling Price with
        #                 ↳Trend Line')
        # fig.add_traces(px.line(grouped_data, x='No_of_Years', y='Selling_Price').data)
        # fig.show()
```

```
[205]: grouped_data = data.groupby('No_of_Years', as_index=False)['Selling_Price'].
        ↳mean().round(2)
        plt.figure(figsize=(10, 6))
        sns.lineplot(data=grouped_data, x='No_of_Years', y='Selling_Price',
        ↳color='blue', label='Data line')
        sns.scatterplot(data=grouped_data, x='No_of_Years', y='Selling_Price',
        ↳color='red', label='Data Points')

        plt.title('Scatter Plot of No of Years vs Selling Price with Trend Line')
        plt.xlabel('No of Years')
```

```
plt.ylabel('Selling Price')
plt.legend()
plt.show()
```




```
[206]: # grouped_data = data.groupby('Year', as_index=False)['Selling_Price'].mean().
        ↪round(2)
        # fig = px.scatter(grouped_data, x='Year', y='Selling_Price',
        #                   title='Scatter plot of Years vs Selling Price with Trend
        ↪Line')
        # fig.add_traces(px.line(grouped_data, x='Year', y='Selling_Price').data)
        # fig.show()
```

```
[207]: grouped_data = data.groupby('Year', as_index=False)['Selling_Price'].mean().
        ↪round(2)
        plt.figure(figsize=(10, 6))
        sns.lineplot(data=grouped_data, x='Year', y='Selling_Price', color='blue',
        ↪label='Data line')
        sns.scatterplot(data=grouped_data, x='Year', y='Selling_Price', color='red',
        ↪label='Data Points')


        plt.title('Scatter Plot of Years vs Selling Price with Trend Line')
        plt.xlabel('Years')
```

```
plt.ylabel('Selling Price')
plt.legend()
plt.show()
```



2- Numerical vs. Categorical (Box Plot)

```
[208]: # Distribution of Selling_Price across different Fuel_Types
plt.figure(figsize=(25, 18))
for i,col in enumerate(categorical_columns):
    plt.subplot(2,len(categorical_columns),i+1)
    sns.boxplot(x=data[col], y='Selling_Price', data=data)
    plt.title(f'Box Plot of Selling_Price by {col}')
    plt.xlabel('Fuel_Type')
    plt.tight_layout()
```




3- Categorical vs. Categorical (Cross-tabulation - Stacked Bar Plot)

```
[209]: pd.crosstab(data['Fuel_Type'], data['Transmission'])
```

```
[209]: Transmission Automatic Manual
Fuel_Type
CNG           0           2
Diesel        11          47
Petrol        28         211
```

displays the frequency distribution of the variables Fuel_Type and Transmission

```
[210]: pd.crosstab(data['Fuel_Type'], data['Transmission']).plot(kind='bar',
    ↪stacked=True)
plt.title('Stacked Bar Plot of Fuel_Type and Transmission')
plt.xlabel('Fuel_Type')
plt.ylabel('Count')
plt.show()
```




10 Multivariate Analysis

(Pair Plot - Heatmap)

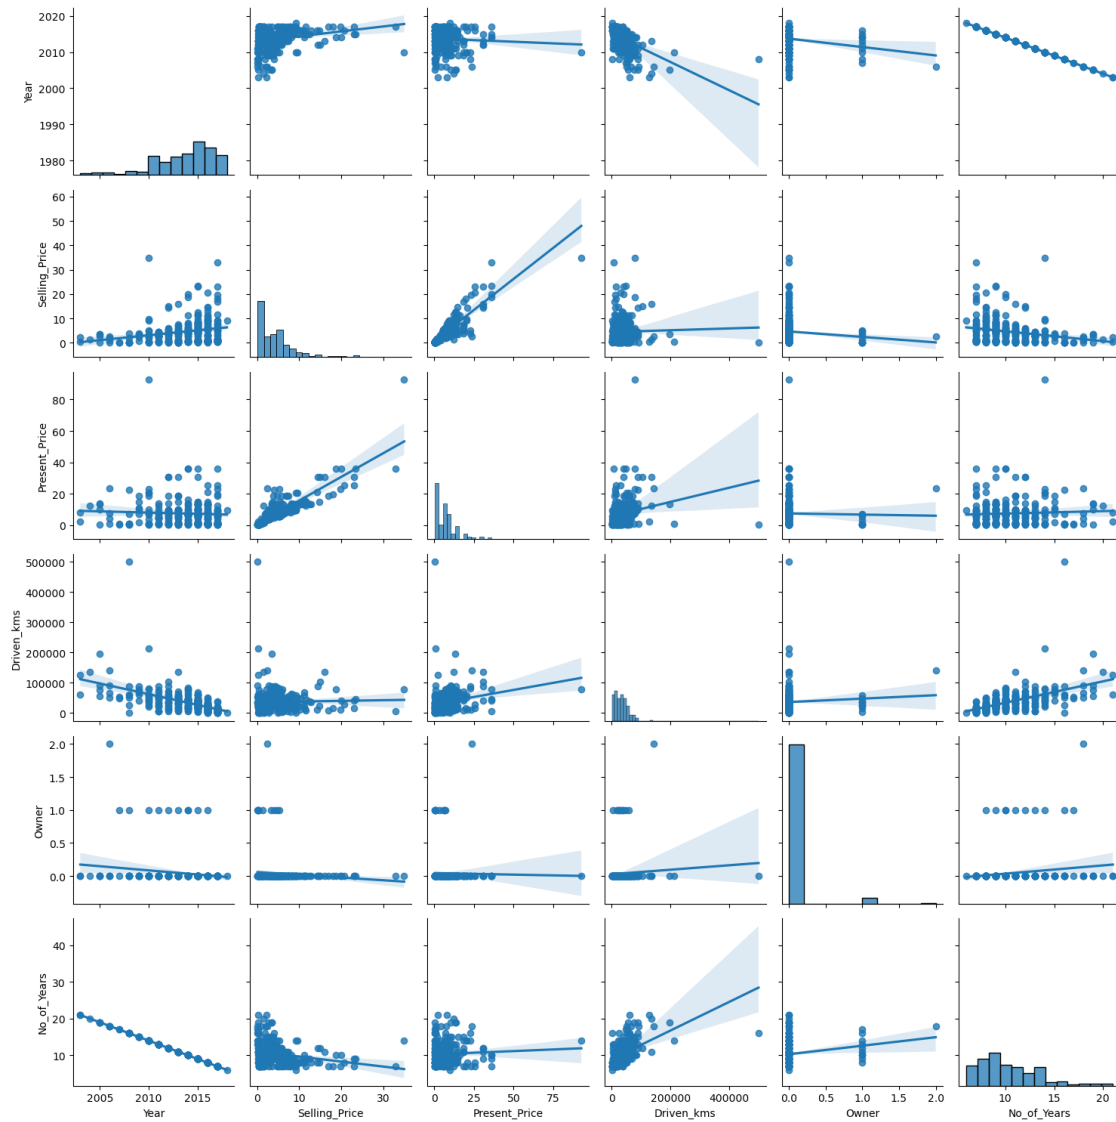
```
[211]: sns.heatmap(correlation_matrix.corr(),annot=True,cmap='coolwarm')  
plt.title('Correlation between effictif features')
```

```
[211]: Text(0.5, 1.0, 'Correlation between effictif features')
```

Strong correlation between Selling_price & Present_price

```
[34]: sns.pairplot(data.select_dtypes('number'),kind='reg')
plt.tight_layout()
```



11 Outliers

```
[212]: Q1 = data['Present_Price'].quantile(0.25)
Q3 = data['Present_Price'].quantile(0.75)
IQR = Q3 - Q1

lower_bound = Q1 - IQR*1.5
upper_bound = Q3 + IQR*1.5

# outliers = New_data[(New_data['Present_Price'] > upper_bound) |
# ↪(New_data['Present_Price'] < lower_bound)]
outliers = data[data.Present_Price > 25.39]
```

```
outliers_count = outliers.count()
outliers_count
```

```
[212]: Year          9
      Selling_Price  9
      Present_Price  9
      Driven_kms    9
      Fuel_Type     9
      Selling_type   9
      Transmission  9
      Owner         9
      No_of_Years   9
      dtype: int64
```

```
[213]: outliers = data[data.Present_Price > 25.39]
      data_clean = data[~(data.Present_Price > 25.39)]
      data_clean
```

```
[213]:      Year  Selling_Price  Present_Price  Driven_kms  Fuel_Type  Selling_type  \
0    2014           3.35           5.59        27000    Petrol      Dealer
1    2013           4.75           9.54        43000    Diesel      Dealer
2    2017           7.25           9.85         6900    Petrol      Dealer
3    2011           2.85           4.15         5200    Petrol      Dealer
4    2014           4.60           6.87        42450    Diesel      Dealer
..    ...           ...           ...           ...           ...           ...
296  2016           9.50          11.60        33988    Diesel      Dealer
297  2015           4.00           5.90        60000    Petrol      Dealer
298  2009           3.35          11.00        87934    Petrol      Dealer
299  2017          11.50          12.50         9000    Diesel      Dealer
300  2016           5.30           5.90         5464    Petrol      Dealer
```

```
      Transmission  Owner  No_of_Years
0      Manual      0      10
1      Manual      0      11
2      Manual      0       7
3      Manual      0      13
4      Manual      0      10
..      ...      ...      ...
296     Manual      0       8
297     Manual      0       9
298     Manual      0      15
299     Manual      0       7
300     Manual      0       8
```

```
[290 rows x 9 columns]
```

```
[214]: outliers.info()
```


```
<class 'pandas.core.frame.DataFrame'>
Index: 9 entries, 50 to 86
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Year             9 non-null     int64
1   Selling_Price    9 non-null     float64
2   Present_Price    9 non-null     float64
3   Driven_kms       9 non-null     int64
4   Fuel_Type        9 non-null     object
5   Selling_type     9 non-null     object
6   Transmission     9 non-null     object
7   Owner            9 non-null     int64
8   No_of_Years      9 non-null     int64
dtypes: float64(2), int64(4), object(3)
memory usage: 720.0+ bytes
```

```
[215]: data_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 290 entries, 0 to 300
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Year             290 non-null    int64
1   Selling_Price    290 non-null    float64
2   Present_Price    290 non-null    float64
3   Driven_kms       290 non-null    int64
4   Fuel_Type        290 non-null    object
5   Selling_type     290 non-null    object
6   Transmission     290 non-null    object
7   Owner            290 non-null    int64
8   No_of_Years      290 non-null    int64
dtypes: float64(2), int64(4), object(3)
memory usage: 22.7+ KB
```

```
[216]: # px.box(data_frame=data_clean,x='Present_Price')
plt.figure(figsize=(10, 6))
sns.boxplot(data=data_clean,x='Present_Price')
```

```
[216]: <Axes: xlabel='Present_Price'>
```



12 Normality

- **Box-Cox Transformation:**

- **Purpose:** Normalize data, reduce skewness, and stabilize variance.
- **Use:** When data is highly skewed and positive. Ideal for preparing data for models that assume normality.

- **Square Root Transformation:**

- **Purpose:** Reduce moderate skewness and stabilize variance.
- **Use:** For non-negative data, especially count data, where a simpler transformation is sufficient.

Both transformations are used to make data more suitable for modeling by addressing issues like skewness and variance instability.

```
[218]: from scipy.stats import boxcox

# Box-Cox transformation to both 'Present_Price' and 'Selling_Price'
data_clean.loc[:, 'Present_Price_transformed'], _ = boxcox(data_clean['Present_Price'] + 1)
data_clean.loc[:, 'Selling_Price_transformed'], _ = boxcox(data_clean['Selling_Price'] + 1)
#square root for Driven-kms
data_clean.loc[:, 'Driven_kms_sqrt'] = np.sqrt(data_clean['Driven_kms'])
```

Plot the original and transformed columns using kdeplot

```
[219]: plt.figure(figsize=(20, 15))

plt.subplot(3, 2, 1)
sns.kdeplot(data=data_clean, x='Present_Price')
plt.title('Present_Price')

plt.subplot(3, 2, 2)
sns.kdeplot(data=data_clean, x='Present_Price_transformed')
plt.title('Box-Cox Transformed Present_Price')

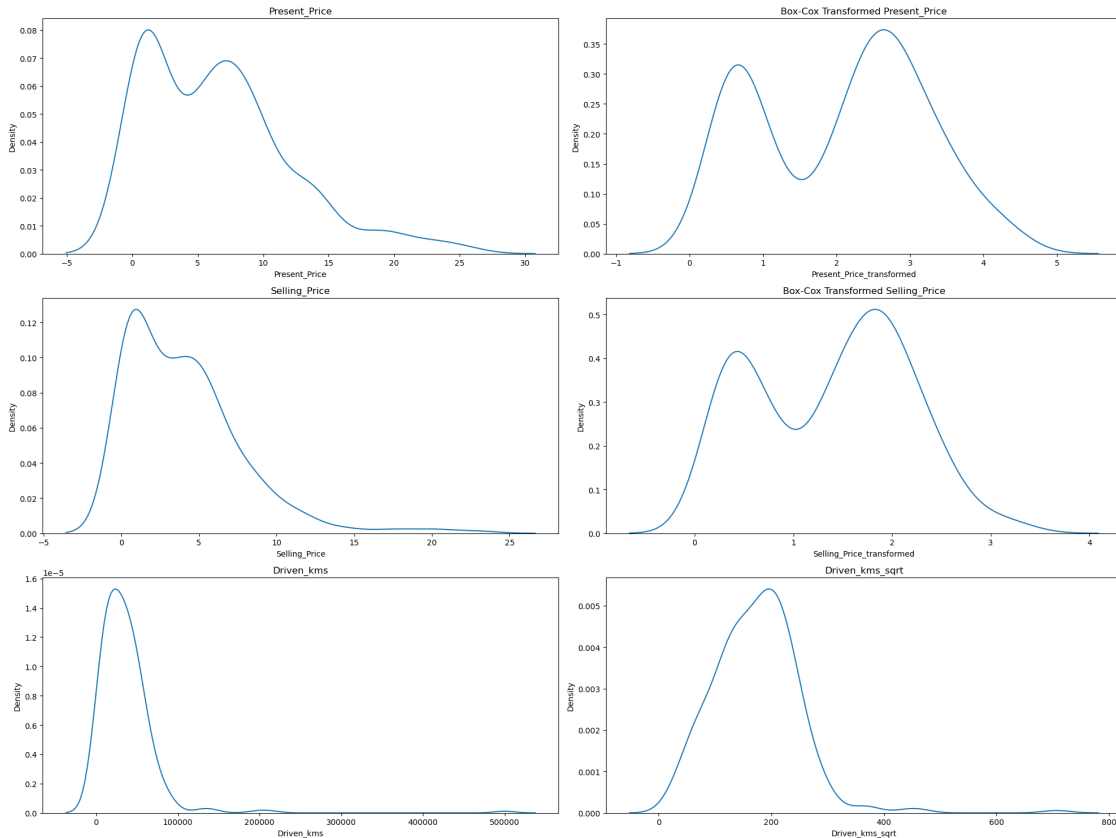
plt.subplot(3, 2, 3)
sns.kdeplot(data=data_clean, x='Selling_Price')
plt.title('Selling_Price')

plt.subplot(3, 2, 4)
sns.kdeplot(data=data_clean, x='Selling_Price_transformed')
plt.title('Box-Cox Transformed Selling_Price')

plt.subplot(3, 2, 5)
sns.kdeplot(data=data_clean, x='Driven_kms')
plt.title('Driven_kms')

plt.subplot(3, 2, 6)
sns.kdeplot(data=data_clean, x='Driven_kms_sqrt')
plt.title('Driven_kms_sqrt')

plt.tight_layout()
plt.show()
```



13 Kurtosis

Kurtosis measures how much the tails of a data distribution differ from a normal distribution.

- **High Kurtosis:** Tails are heavy, meaning there are more extreme values (outliers) than usual.
- **Low Kurtosis:** Tails are light, meaning there are fewer extreme values than usual.
- **Normal Kurtosis:** Tails are similar to those of a normal distribution, with a balanced number of outliers.

```
[220]: from scipy.stats import kurtosis
kurtosis_present_price = kurtosis(data_clean['Present_Price'])
kurtosis_selling_price = kurtosis(data_clean['Selling_Price'])
kurtosis_Driven_kms = kurtosis(data_clean['Driven_kms'])

print(f"Kurtosis of Present_Price: {kurtosis_present_price}")
print(f"Kurtosis of Selling_Price: {kurtosis_selling_price}")
print(f"Kurtosis of Driven_kms: {kurtosis_Driven_kms}")
```

Kurtosis of Present_Price: 0.7768684598078957

Kurtosis of Selling_Price: 4.182781604796288

Kurtosis of Driven_kms: 72.34036114182679

Before Transformation:

1. **Present_Price:**
 - **Kurtosis:** 0.7769
 - **Interpretation: Mesokurtic.** The distribution is close to normal with moderate peak and tails.
2. **Selling_Price:**
 - **Kurtosis:** 4.1828
 - **Interpretation: Leptokurtic.** This indicates a distribution with a higher peak and heavier tails, suggesting a greater presence of extreme values.
3. **Driven_kms:**
 - **Kurtosis:** 72.3404
 - **Interpretation: Highly Leptokurtic.** This extreme kurtosis shows a distribution with a very sharp peak and extremely heavy tails, indicating many extreme values and outliers.

```
[221]: kurtosis_present_price_transformed =  
↳ kurtosis(data_clean['Present_Price_transformed'])  
kurtosis_selling_price_transformed =  
↳ kurtosis(data_clean['Selling_Price_transformed'])  
kurtosis_Driven_kms_sqrt = kurtosis(data_clean['Driven_kms_sqrt'])  
  
print(f"Kurtosis of Present_Price_transformed: {  
↳ kurtosis_present_price_transformed}")  
print(f"Kurtosis of Selling_Price_transformed: {  
↳ kurtosis_selling_price_transformed}")  
print(f"Kurtosis of Driven_kms_sqrt: {kurtosis_Driven_kms_sqrt}")
```

Kurtosis of Present_Price_transformed: -1.1828042960228728

Kurtosis of Selling_Price_transformed: -1.035266774316057

Kurtosis of Driven_kms_sqrt: 7.824823450468216

After Transformation:

1. **Present_Price_transformed:**
 - **Kurtosis:** -1.1828
 - **Interpretation: Platykurtic.** The transformed distribution is flatter than normal, with fewer extreme values compared to the original Present_Price.
2. **Selling_Price_transformed:**
 - **Kurtosis:** -1.0353
 - **Interpretation: Platykurtic.** The transformed Selling_Price also shows a flatter distribution with fewer extreme values, improving the normality of the data.
3. **Driven_kms_sqrt:**
 - **Kurtosis:** 7.8248
 - **Interpretation: Leptokurtic.** Although the transformation made the distribution less extreme compared to the original, it still has heavy tails and a higher peak, indicating a significant presence of extreme values.

Conclusion: After Transformation: The transformed data generally shows improvements in terms of reducing extreme values, making it closer to normal distribution. This is usually beneficial for model performance

```
[222]: data_clean.drop(columns=['Driven_kms', 'Selling_Price', 'Present_Price'],  
    ↪ inplace=True)
```

C:\Users\lenovo\AppData\Local\Temp\ipykernel_24308\361000944.py:1:
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
[223]: data_clean.columns
```

```
[223]: Index(['Year', 'Fuel_Type', 'Selling_type', 'Transmission', 'Owner',  
        'No_of_Years', 'Present_Price_transformed', 'Selling_Price_transformed',  
        'Driven_kms_sqrt'],  
        dtype='object')
```

```
[224]: from sklearn.preprocessing import LabelEncoder  
  
data_clean = data_clean.copy()  
le = LabelEncoder()  
  
# LabelEncoder to each column  
data_clean['Fuel_Type'] = le.fit_transform(data_clean['Fuel_Type'])  
data_clean['Selling_type'] = le.fit_transform(data_clean['Selling_type'])  
data_clean['Transmission'] = le.fit_transform(data_clean['Transmission'])  
  
data_clean.head()
```

```
[224]:
```

	Year	Fuel_Type	Selling_type	Transmission	Owner	No_of_Years	\
0	2014	2	0	1	0	10	
1	2013	1	0	1	0	11	
2	2017	2	0	1	0	7	
3	2011	2	0	1	0	13	
4	2014	1	0	1	0	10	

	Present_Price_transformed	Selling_Price_transformed	Driven_kms_sqrt
0	2.25	1.50	164.32
1	2.94	1.80	207.36
2	2.98	2.18	83.07
3	1.91	1.38	72.11

4	2.50	1.77	206.03
---	------	------	--------

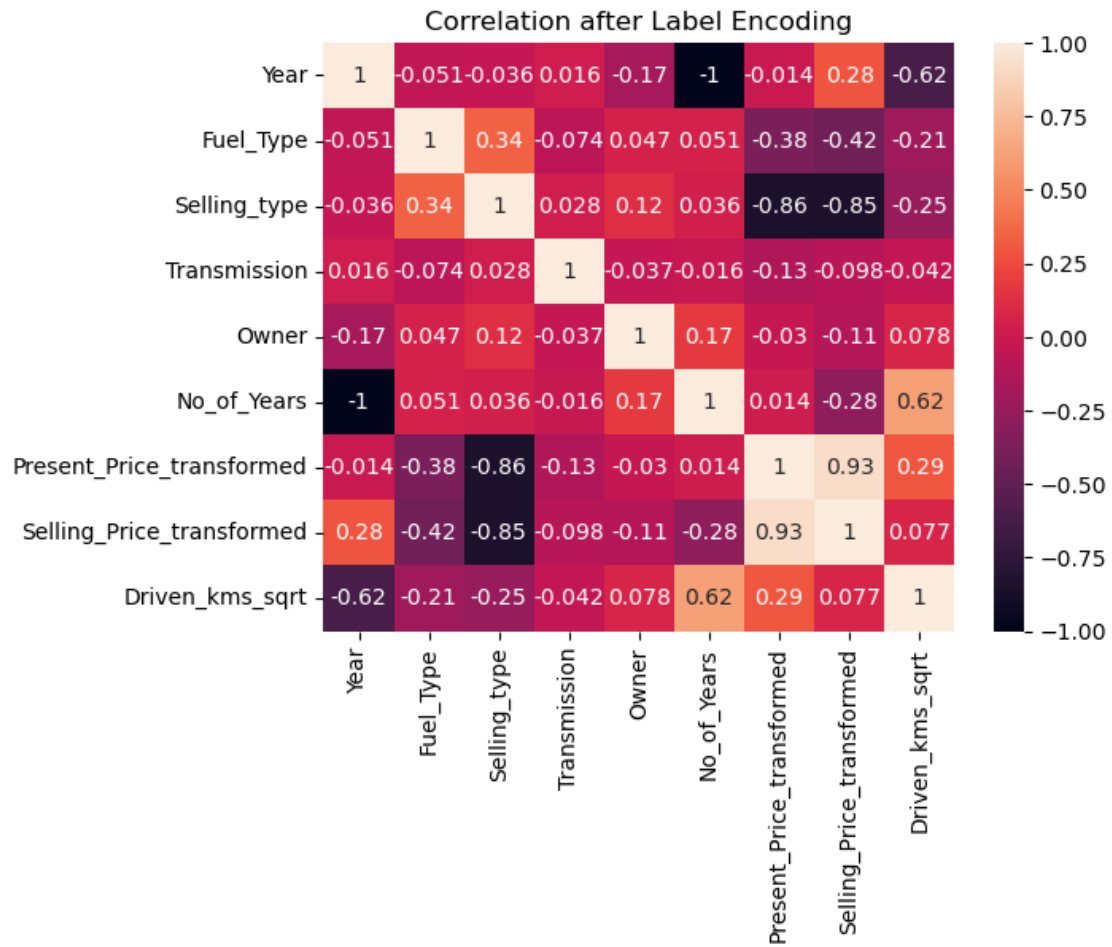
```
[225]: data_clean.dtypes
```

```
[225]: Year                int64
      Fuel_Type          int32
      Selling_type       int32
      Transmission       int32
      Owner              int64
      No_of_Years        int64
      Present_Price_transformed float64
      Selling_Price_transformed float64
      Driven_kms_sqrt     float64
      dtype: object
```

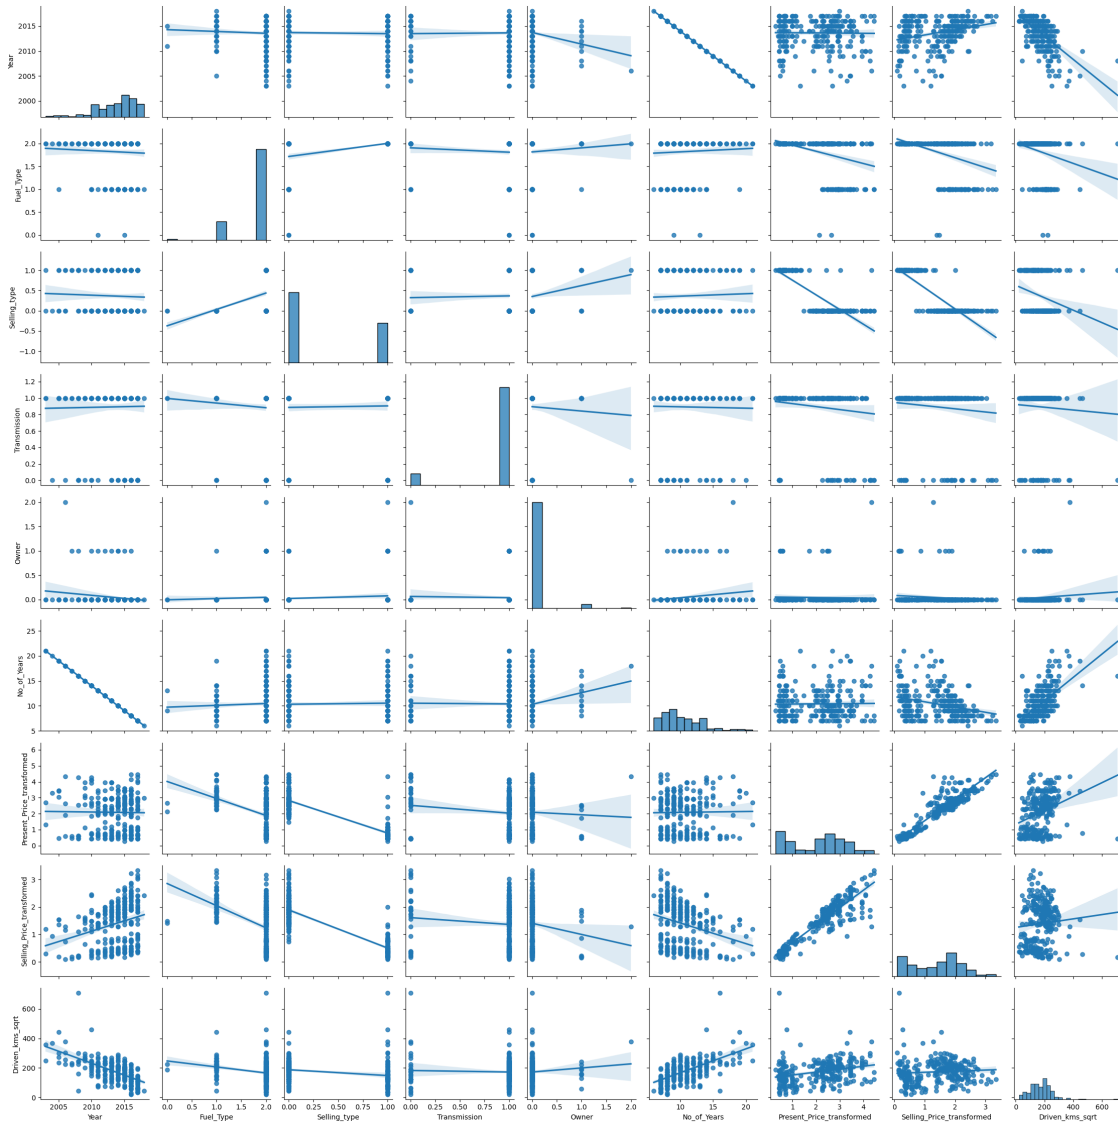
14 Correlations__After__Encoding

```
[226]: sns.heatmap(data_clean.corr(),annot=True)
      plt.title('Correlation after Label Encoding')
```

```
[226]: Text(0.5, 1.0, 'Correlation after Label Encoding')
```



```
[227]: sns.pairplot(data_clean,kind='reg')
plt.tight_layout()
```



15 Multivariate Regression Analysis

Multivariate Regression: Analyze the relationship between one dependent variable and multiple independent variables.

```
[228]: import statsmodels.api as sm
x = data_clean.drop(['Selling_Price_transformed'],axis=1)
y = data_clean['Selling_Price_transformed']
```

```
[229]: x = sm.add_constant(x)
```

```
[230]: model = sm.OLS(y, x).fit()
```

```
[231]: x.columns
```

```
[231]: Index(['const', 'Year', 'Fuel_Type', 'Selling_type', 'Transmission', 'Owner',  
        'No_of_Years', 'Present_Price_transformed', 'Driven_kms_sqrt'],  
        dtype='object')
```

16 Assumptions in Ordinary Least Squares (OLS) regression

1- Normality of Residuals

```
[232]: residuals = model.resid
```

- Anderson-Darling

```
[233]: # Anderson-Darling test on residuals  
from scipy.stats import anderson  
  
result_residuals = anderson(residuals, dist='norm')  
print("Results for OLS Residuals:")  
print(f"Anderson-Darling Test Statistic: {round(result_residuals.statistic, 4)}")  
print(f"Critical Values: {result_residuals.critical_values}")  
print(f"Significance Levels: {result_residuals.significance_level}")
```

Results for OLS Residuals:

Anderson-Darling Test Statistic: 0.5968

Critical Values: [0.568 0.647 0.777 0.906 1.077]

Significance Levels: [15. 10. 5. 2.5 1.]

The Anderson-Darling test statistic for the OLS residuals is 0.5968.

- **At the 15% significance level:** The test statistic (0.5968) is slightly above the critical value (0.568). This indicates a minor deviation from normality, but it's minimal.
- **At the 10% significance level:** The test statistic (0.5968) is below the critical value (0.647). Thus, we do not reject the null hypothesis, suggesting the residuals are consistent with normality at this level.
- **At the 5%, 2.5%, and 1% significance levels:** The test statistic is below the critical values (0.777, 0.906, and 1.077, respectively). Therefore, we do not reject the null hypothesis at these levels, meaning the residuals are close to a normal distribution.
- Shapiro-Wilk

```
[234]: from scipy.stats import shapiro  
# Perform Shapiro-Wilk test on residuals  
stat, p_value = shapiro(residuals)  
print("Results for OLS Residuals:")  
print(f"Shapiro-Wilk Test Statistic: {round(stat, 4)}")  
print(f"p-value: {round(p_value, 4)}")
```

```
# Interpretation
if p_value > 0.05:
    print("Conclusion: The residuals are likely normally distributed (fail to
    ↪reject the null hypothesis)")
else:
    print("Conclusion: The residuals are not normally distributed (reject the
    ↪null hypothesis)")
```

Results for OLS Residuals:

Shapiro-Wilk Test Statistic: 0.9908


p-value: 0.0672

Conclusion: The residuals are likely normally distributed (fail to reject the null hypothesis)

- Histogram of Residuals


```
[235]: sns.kdeplot(residuals)
```

```
[235]: <Axes: ylabel='Density'>
```



- QQ Plot

```
[236]: stats.probplot(residuals, dist="norm", plot=plt)
plt.title("QQ Plot of OLS Residuals")
plt.show()
```



2-Independence

The Durbin-Watson Statistic (DW) ranges from 0 to 4:

- **A value of 2** indicates no autocorrelation in the residuals.
- **Values closer to 0** suggest positive autocorrelation, meaning the residuals are positively correlated.
- **Values closer to 4** indicate negative autocorrelation, meaning the residuals are negatively correlated.

```
[237]: dw = sm.stats.durbin_watson(residuals)
dw
```

```
[237]: 1.7698085280996163
```

1.7698, it is close to 2, suggesting that there is no strong evidence of autocorrelation in the residuals of your regression model. This is generally a good sign, indicating that your residuals are approximately uncorrelated.

[238]: `model.summary()`

[238]:

Dep. Variable:	Selling_Price_transformed	R-squared:	0.956
Model:	OLS	Adj. R-squared:	0.955
Method:	Least Squares	F-statistic:	878.6
Date:	Thu, 29 Aug 2024	Prob (F-statistic):	2.12e-187
Time:	23:04:28	Log-Likelihood:	113.29
No. Observations:	290	AIC:	-210.6
Df Residuals:	282	BIC:	-181.2
Df Model:	7		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-3.214e-05	2.21e-06	-14.553	0.000	-3.65e-05	-2.78e-05
Year	0.0007	4.5e-05	15.228	0.000	0.001	0.001
Fuel_Type	-0.1268	0.027	-4.713	0.000	-0.180	-0.074
Selling_type	-0.2638	0.041	-6.419	0.000	-0.345	-0.183
Transmission	-0.0042	0.033	-0.130	0.897	-0.068	0.060
Owner	-0.0691	0.047	-1.482	0.139	-0.161	0.023
No_of_Years	-0.0657	0.004	-14.682	0.000	-0.075	-0.057
Present_Price_transformed	0.5361	0.018	30.085	0.000	0.501	0.571
Driven_kms_sqrt	-0.0005	0.000	-3.017	0.003	-0.001	-0.000

Omnibus:	4.910	Durbin-Watson:	1.770
Prob(Omnibus):	0.086	Jarque-Bera (JB):	5.136
Skew:	-0.201	Prob(JB):	0.0767
Kurtosis:	3.513	Cond. No.	7.31e+18

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 2.22e-29. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

R-squared = 0.95 that means i can dependent on these features to determine car price bec squared_correlation is high

17 Remove Insignificant Predictors

p-value > 0.05 (Transmission - Owner)

```
[239]: x=x.drop(['Transmission'],axis=1)
model=sm.OLS(y,x).fit()
model.summary()
```

[239]:

Dep. Variable:	Selling_Price_transformed	R-squared:	0.956
Model:	OLS	Adj. R-squared:	0.955
Method:	Least Squares	F-statistic:	1029.
Date:	Thu, 29 Aug 2024	Prob (F-statistic):	6.38e-189
Time:	23:04:30	Log-Likelihood:	113.28
No. Observations:	290	AIC:	-212.6
Df Residuals:	283	BIC:	-186.9
Df Model:	6		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-3.216e-05	2.2e-06	-14.594	0.000	-3.65e-05	-2.78e-05
Year	0.0007	3.87e-05	17.624	0.000	0.001	0.001
Fuel_Type	-0.1264	0.027	-4.747	0.000	-0.179	-0.074
Selling_type	-0.2629	0.040	-6.497	0.000	-0.343	-0.183
Owner	-0.0690	0.047	-1.484	0.139	-0.161	0.023
No_of_Years	-0.0658	0.004	-14.720	0.000	-0.075	-0.057
Present_Price_transformed	0.5367	0.017	30.964	0.000	0.503	0.571
Driven_kms_sqrt	-0.0005	0.000	-3.019	0.003	-0.001	-0.000

Omnibus:	4.619	Durbin-Watson:	1.769
Prob(Omnibus):	0.099	Jarque-Bera (JB):	4.786
Skew:	-0.192	Prob(JB):	0.0914
Kurtosis:	3.498	Cond. No.	7.31e+18

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The smallest eigenvalue is 2.22e-29. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
[240]: x=x.drop(['Owner'],axis=1)
model=sm.OLS(y,x).fit()
model.summary()
```

```
[240]:
```

Dep. Variable:	Selling_Price_transformed	R-squared:	0.956
Model:	OLS	Adj. R-squared:	0.955
Method:	Least Squares	F-statistic:	1229.
Date:	Thu, 29 Aug 2024	Prob (F-statistic):	5.18e-190
Time:	23:04:31	Log-Likelihood:	112.16
No. Observations:	290	AIC:	-212.3
Df Residuals:	284	BIC:	-190.3
Df Model:	5		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-3.256e-05	2.19e-06	-14.862	0.000	-3.69e-05	-2.82e-05
Year	0.0007	3.83e-05	18.030	0.000	0.001	0.001
Fuel_Type	-0.1272	0.027	-4.770	0.000	-0.180	-0.075
Selling_type	-0.2732	0.040	-6.840	0.000	-0.352	-0.195
No_of_Years	-0.0666	0.004	-14.992	0.000	-0.075	-0.058
Present_Price_transformed	0.5332	0.017	30.985	0.000	0.499	0.567
Driven_kms_sqrt	-0.0005	0.000	-2.997	0.003	-0.001	-0.000
<hr/>						
Omnibus:	5.102	Durbin-Watson:		1.791		
Prob(Omnibus):	0.078	Jarque-Bera (JB):		5.571		
Skew:	-0.189	Prob(JB):		0.0617		
Kurtosis:	3.564	Cond. No.		7.31e+18		

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 2.22e-29. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

18 OLS_Parameters

```
[241]: print("R-squared value:")
print(model.rsquared)

print("\nMSE value:")
print(model.mse_resid)

print("\nRSE value:")
rse = np.sqrt(model.mse_resid)
print(rse)

print("\nFitted Values:")
print(model.fittedvalues.head())

print("\nResiduals:")
print(model.resid.head())
```

R-squared value:
0.9558122939762361

MSE value:
0.02758526181987797

RSE value:
0.1660881146255745

Fitted Values:
0 1.58

```
1    1.99
2    2.22
3    1.25
4    1.82
dtype: float64
```

```
Residuals:
0    -0.08
1    -0.19
2    -0.04
3     0.13
4    -0.05
dtype: float64
```

```
[242]: x.columns
```

```
[242]: Index(['const', 'Year', 'Fuel_Type', 'Selling_type', 'No_of_Years',
          'Present_Price_transformed', 'Driven_kms_sqrt'],
          dtype='object')
```

19 Car Price Prediction

```
[243]: from sklearn.model_selection import train_test_split
       from sklearn.preprocessing import StandardScaler
       from sklearn.metrics import mean_squared_error, r2_score
```

```
[244]: y = data_clean['Selling_Price_transformed']
       X = data_clean[['Year', 'Fuel_Type', 'Selling_type', 'No_of_Years',
          'Present_Price_transformed', 'Driven_kms_sqrt']]
```

```
[245]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
       ↪ 2, random_state=100)
```

** Scaling Data

```
[246]: scaler = StandardScaler()
       x_train_scaled = scaler.fit_transform(X_train)
       x_test_scaled = scaler.transform(X_test)
```

1- Gradient Boosting

```
[247]: from sklearn.ensemble import GradientBoostingRegressor
       model = GradientBoostingRegressor(n_estimators=100, learning_rate=0.1,
       ↪ max_depth=3, random_state=100)
```

```
[248]: model.fit(x_train_scaled, y_train)
```

```
[248]: GradientBoostingRegressor(random_state=100)
```

```
[249]: y_pred=model.predict(x_test_scaled)
```

```
[250]: mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f"Mean Squared Error: {mse}")
print(f"R-squared: {r2}")
```

Mean Squared Error: 0.01275779694321093
R-squared: 0.9746054223909847

```
[251]: gradient_data=pd.DataFrame({'Actual':y_test , 'Predicted':y_pred})
gradient_data.reset_index(inplace=True,drop=True)
gradient_data.head()
```

```
[251]:
```

	Actual	Predicted
0	2.41	2.35
1	0.93	1.40
2	0.47	0.37
3	0.22	0.28
4	1.72	1.84

2- Decision Tree

```
[252]: from sklearn.tree import DecisionTreeRegressor
model = DecisionTreeRegressor(max_depth=None, random_state=100)
```

```
[253]: model.fit(x_train_scaled, y_train)
```

```
[253]: DecisionTreeRegressor(random_state=100)
```

```
[254]: y_pred=model.predict(x_test_scaled)
```

```
[255]: mse=mean_squared_error(y_pred,y_test)
r2=r2_score(y_pred,y_test)
print(f'Mean Squared Error is : {mse}')
print(f'R-squared is : {r2}')
```

Mean Squared Error is : 0.023709525712562646
R-squared is : 0.9564019812412028

```
[256]: desicion_data=pd.DataFrame({'Actaul':y_test , 'Predicted':y_pred})
desicion_data.reset_index(inplace=True,drop=True)
desicion_data.head()
```

```
[256]:
```

	Actaul	Predicted
0	2.41	2.45
1	0.93	1.54
2	0.47	0.34
3	0.22	0.18

4 1.72 1.67

3- Lasso

```
[257]: from sklearn.linear_model import Lasso
lasso = Lasso(alpha=0.1) # You can adjust the alpha parameter for
↳ regularization strength
```

```
[258]: lasso.fit(x_train_scaled, y_train)
```

```
[258]: Lasso(alpha=0.1)
```

```
[259]: y_pred = model.predict(x_test_scaled)

# Evaluate model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse}")
print(f"R-squared: {r2}")
```

Mean Squared Error: 0.023709525712562646

R-squared: 0.9528058493593581

```
[260]: lasso_data=pd.DataFrame({'Actaul':y_test , 'Predicted':y_pred})
lasso_data.reset_index(inplace=True,drop=True)
lasso_data.head()
```

```
[260]:      Actaul   Predicted
0      2.41          2.45
1      0.93          1.54
2      0.47          0.34
3      0.22          0.18
4      1.72          1.67
```

20 Feature selection using Lasso

```
[261]: lasso_coefficients = pd.DataFrame({
        'Feature': X.columns,
        'Coefficient': lasso.coef_
    })

# Select important features (non-zero coefficients)
important_features = lasso_coefficients[lasso_coefficients['Coefficient'] != 0]

print("Important features selected by Lasso:")
print(important_features)
```

Important features selected by Lasso:

	Feature	Coefficient
0	Year	0.10
2	Selling_type	-0.11
3	No_of_Years	-0.02
4	Present_Price_transformed	0.54


Year - Present_Price_transformed is the most important

- **Selected Features:** Those with non-zero coefficients after applying Lasso.
- **Excluded Features:** Those with zero or less coefficients

```
[262]: import matplotlib.pyplot as plt
import seaborn as sns

# Plot the important features selected by Lasso
plt.figure(figsize=(10, 6))
sns.barplot(y='Coefficient', x='Feature', data=important_features.
    ↪sort_values(by='Coefficient', ascending=False))

plt.title('Feature Importance by Lasso Regression')
plt.xlabel('Coefficient Value')
plt.ylabel('Feature')
plt.show()
```



4- Linear Regression

```
[263]: from sklearn.linear_model import LinearRegression
model=LinearRegression()
```

```
[264]: model.fit(X_train,y_train)
```

```
[264]: LinearRegression()
```

```
[ ]: y_pred=model.predict(X_test)
```

```
[265]: mse=mean_squared_error(y_pred,y_test)
r2=r2_score(y_pred,y_test)
print(f'Mean Squared Error is : {mse}')
print(f'R-squared is : {r2}')
```

Mean Squared Error is : 0.023709525712562646
R-squared is : 0.9564019812412028

```
[266]: linear_data=pd.DataFrame({'Actaul':y_test , 'Predicted':y_pred})
linear_data.reset_index(inplace=True,drop=True)
linear_data.head()
```

```
[266]:
```

	Actaul	Predicted
0	2.41	2.45
1	0.93	1.54
2	0.47	0.34
3	0.22	0.18
4	1.72	1.67


21 Plot True Values vs. Predicted Values

```
[267]: plt.figure(figsize=(10, 6))
plt.scatter(gradient_data['Actual'], gradient_data['Predicted'], color='blue',
            ↪label='Predicted vs Actual')

# Line for perfect prediction
min_value = gradient_data['Actual'].min()
max_value = gradient_data['Actual'].max()
plt.plot([min_value, max_value], [min_value, max_value], color='red',
        ↪linestyle='--', label='Perfect Prediction')

plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title('True Values vs. Predicted Values')
plt.legend()

plt.show()
```



```
[268]: results=pd.DataFrame({
        'Model': ['Gradient Boosting', 'Decision Tree', 'Lasso', 'LinearRegression'],
        'Score': [0.97, 0.956, 0.952, 0.96]
    })
result_pd=results.sort_values(by='Score', ascending=False)
result_pd=result_pd.set_index('Score')
result_pd
```

```
[268]:
```

	Model
Score	
0.97	Gradient Boosting
0.96	LinearRegression
0.96	Decision Tree
0.95	Lasso

```
[ ]:
```

```
[ ]:
```